
Independent Verification Paths Are Not Independent: A Case Study of Common-Mode Failure in a Satellite Catalogue Pipeline

Fabio Rovai
The Tesseract Academy
London, United Kingdom
fabio@thetesseractacademy.com

Abstract

A common safeguard for a data pipeline is redundant computation: derive each published number by two routes built on different technology and refuse to exit when they disagree. We report one such gate failing, in a cross-catalogue integrity study of two open registers of Earth-orbiting objects. A gate comparing a set-based Python path with SPARQL queries over the emitted RDF graph printed ALL CROSS-CHECKS AGREE on seven counts. Three were wrong, one overstated more than fourfold (932 against 220). Both paths imported the same constants, which encoded a misreading of the source’s status vocabulary, so the error was common-mode and the gate could not see it. We give the mechanism, an object-level ledger reconciling every figure, and three checks that go back to the source’s documentation, measured on the defective code and on its correction. We then checked that correction against each object’s phase history, held in a source file the pipeline never read. The correction was also wrong: 42 of its 261 disagreements are artefacts, and none of our three checks flagged them. Finally, in a controlled replication with three pinned models and tools disabled, 72 of 75 paths generated on request as independent checks computed the defective count, 29 of 30 even when the prompt carried the source’s own definitions of the codes. The evidence is one pipeline and one defect family. Within it, redundancy verified implementation, and the errors that reached publication were errors of meaning.

1 Introduction

A pipeline that produces more numbers than its authors can trace by hand has to check itself, and one standard way to do it is redundancy: compute every quantity that will be published by two routes built on different technology, compare them, and refuse to exit successfully unless they agree. The check is cheap, mechanical, and needs no ground truth, which is why it appeals wherever findings are generated faster than they can be audited, including in systems that propose and test hypotheses automatically [Lu et al., 2024]. The same intuition, that agreement between routes stands in for a correctness signal, underwrites self-consistency decoding [Wang et al., 2023] and the use of agreement between generated code and generated tests to rank programs [Chen et al., 2023].

This paper is a case study of one such gate failing, in a pipeline we built. It is a single incident, and we do not claim that it tells us how often redundancy gates fail. What it can do is show one failure completely: the gate, the wrong and corrected numbers, the definition that caused the error, and what did and did not catch it.

The pipeline compares two independently maintained public catalogues of artificial Earth-orbiting objects and emits an RDF graph of typed defect assertions. A dual-computation gate counts each

defect class twice, once set-based in Python from the source files and once by SPARQL over the emitted graph, and exits non-zero on any disagreement. On the run whose numbers were published it printed ALL CROSS-CHECKS AGREE. Three of the seven counts it compared were wrong, and the largest was overstated more than fourfold. Both paths imported the same classification constants, and those constants encoded a misreading of the source’s vocabulary of flight-phase status codes. Two correct implementations of one wrong reading agree perfectly, so the error was common-mode and the gate was silent by construction. The multiversion programming literature established long ago that separately developed versions of a program do not fail independently [Knight and Leveson, 1986, Eckhardt and Lee, 1985, Littlewood and Miller, 1989]; our case is a data-pipeline instance of that result in which the shared artifact is not a specification but an executable set of constants.

The contributions are four, each bounded to this pipeline. First, the mechanism, with an object-level ledger that reconciles every figure across versions of the count (Sections 3 and 4, Appendix A). Second, three checks that go back to the source’s own documentation, each implemented and run on the defective code and on its correction (Sections 5 and 7). Third, a second-order failure. Checking our correction against every object’s phase history showed that it misread the same vocabulary in a different way, and that none of our three checks had noticed (Section 6). Fourth, an experiment on whether a code-generating model asked for an independent verification path produces one (Section 8).

Redundancy remains worth having; it verifies implementation. The errors that reached publication here were errors of meaning, which two paths built on shared definitions cannot test.

2 Setting

2.1 Two registers of one population

Artificial objects in Earth orbit are catalogued more than once. We used two sources, both openly retrievable without credentials, both retrieved on 18 August 2026. CelesTrak’s satellite catalogue [CelesTrak, 2026] carried 70,292 objects. The General Catalog of Artificial Space Objects (GCAT) [McDowell, 2020], maintained independently and published under CC BY 4.0, is split across many files. The pipeline read four of them: a main catalogue of 69,999 records covering 69,391 numbered objects, an auxiliary catalogue of 11,962 records, a catalogue of 1,802 objects that failed to orbit, and a catalogue of 354 numbered objects above 100,000. Two registers describing one population is a useful setting because neither has to be treated as ground truth: where they disagree, the disagreement itself is the finding.

2.2 The pipeline and its gate

The pipeline reconciles objects by catalogue number, emits an RDF graph in which every defect is a typed node carrying its evidence, and validates that graph against three layers of SHACL shapes [Knublauch and Kontokostas, 2017], one shape per defect class. On the published run it gated seven defect classes, among them entries one register says correspond to no real object while the other still carries them, and objects on which the two registers disagree about whether the object is still in orbit.

Verification was deliberately over-provisioned. Three independent checks ran on every build: a SHACL processor over all three shape layers, a second RDF engine written in a different language that loaded and validated the same ontology, and a dual-computation gate that recomputed every published count twice. The third is the subject of this paper, and its structure is the obvious one:

```
python_side() reads the source files and returns seven counts. sparql_side() parses
the emitted graph and issues one SELECT (COUNT(DISTINCT ?n)) per defect class.
main() compares them key by key and calls sys.exit(1) if any pair differs.
```

The diversity is real. One path never constructs a graph and the other never sees the source files. One counts Python set members, the other counts distinct SPARQL bindings over a materialised graph of 2,330,660 triples, two orders of magnitude larger than the input. Parsing that graph takes over an hour and roughly 5 GB of memory, so the gate is not run casually, which is part of why its verdict carried weight. On the published run it agreed on all seven counts.

Table 1: Gate output on the published run. Every row agreed between the two paths, and three rows were wrong. The first correction is the value published as a fix; the full-history value is the check in Section 6.

Defect class	Python	SPARQL	Agree	First correction	Full history
PhantomEntry	22	22	yes	22	22
PhantomEntryOnOrbit	1	1	yes	1	1
UndisclosedTrackingLoss	1,104	1,104	yes	1,094	1,094
DispositionDisagreement	932	932	yes	261	220
CoverageGap	900	900	yes	622	622
UnnumberedObject	605	605	yes	605	605
IdentifierCollision	3	3	yes	3	3

3 The failure

Table 1 reproduces the gate’s output on the published run, next to two later values for each count: the first correction, and the value after the check against object histories described in Section 6. The gate exited zero. Every count in this paper is a number of objects present in the catalogues concerned, and Appendix A tracks each disposition object through all three versions.

The three errors have different characters, and only the first is the common-mode failure this paper is about.

Disposition disagreement: 932, then 261, then 220. The count is the number of objects on which the two catalogues contradict each other about whether the object is still in orbit. The 932 comprised 163 objects GCAT records as gone while CelesTrak gives no decay date, and 769 objects CelesTrak records as decayed while GCAT does not record them as gone. Producing the count requires reading GCAT’s status codes, and the pre-correction rule was simple: a status in a set named GONE meant the object was gone, and every other status meant it was still in orbit. That default did most of the damage. Of the 769, 640 were objects whose GCAT record ends in a code that marks a transition rather than a fate, such as docking, attachment, transfer or renaming, and 500 of those were docking or attachment specifically. Another 57 ended in an explosion or a collision. The first correction removed 998 transition-coded objects from the comparison, 640 of which had been counted as disagreements, and reclassified explosion and collision as destruction, giving 261. The ledger in Appendix A accounts for every object that entered or left. Section 6 shows that 261 was itself wrong.

Coverage gap: 900 against 622. The pipeline checked three of the four main catalogue files the source publishes. The fourth was listed on the same index page and never fetched. It holds 354 objects, 278 of which were being counted as absent from a catalogue that contains them.

Undisclosed tracking loss: 1,104 against 1,094. Ten of the 1,104 objects do carry the CelesTrak field whose absence the count is about. Correcting this strengthened the finding. The field is maintained and set on 1,292 objects across the catalogue, ten of them among the 1,104, so the claim moved from “this register cannot record tracking loss” to “this register records tracking loss and does not record it for 1,094 objects.”

None of the three errors is visible to a redundancy check, and only the first is common-mode in the strict sense. What unites them is that each is an error about what the source data means or contains, and the gate compares only what two implementations derive from it.

4 Why the gate could not see it

4.1 The diversity was in the wrong place

The two paths differ in language, data structure and intermediate representation. They do not differ in where their meaning comes from. The Python path imports its classification constants and file loaders from the reconciliation module. The SPARQL path queries a graph built by a script that imports the same constants and loaders from the same module. We computed the overlap from the

source at the published commit: the names each path imports from that module are identical, four vocabularies and two loaders (Appendix D, `shared_ancestor_audit.py`).

```
reconcile.py defines GONE, LEFT_EARTH, ERROR, LOST, load_gcat, load_celestrak
↓
governance_report.py → Python counts build_graph.py → RDF graph → SPARQL counts ↓
```

The gate diversified the last step of the computation and shared the first. Whatever those sets assert about the source vocabulary is asserted identically on both sides, so the agreement the gate reports is evidence that two pieces of code, given the same premises, drew the same conclusion. In one sentence: the gate verified that the Python and SPARQL paths implemented the same misunderstanding identically.

4.2 The defective definition

The pre-correction disposition rule was one line applied to a set literal. An object counted as gone in GCAT if any of its status codes was in

```
GONE = {"R", "D", "L", "LF", "S", "F", "AF", "AS", "AR", "AR IN", "AL", "AL IN", "TX"}
```

and, through `bool(st & GONE) != (n in ctdec)`, every other object counted as still in orbit. The error lived in that default. Codes the set did not name, among them E (exploded), C (collided), DK (docked) and ATT (attached), were silently read as “in orbit”. A rule that made a claim only on explicit membership of GONE or of the in-orbit set, and made no claim otherwise, returns 208 on the frozen data, and all 208 are confirmed by the full-history check of Section 6. The missing codes mattered far less than the decision about what to do with codes nobody had classified.

A smaller defect in the same module shows the pathology from another angle. GCAT marks uncertain dates with a trailing question mark, so a status can arrive as “R?”. The original code compared raw strings and failed to match it, which moved 15 objects in or out of the count. None of these corrections is discoverable by comparing outputs. Each is a correction to a reading of the source.

4.3 Relation to multiversion programming

Independently produced program versions do not fail independently [Knight and Leveson, 1986], and the accompanying theory located the cause in the input rather than the versions: some inputs are hard for everyone, so failures correlate however separately the versions were written [Eckhardt and Lee, 1985, Littlewood and Miller, 1989]. N-version designs were proposed precisely to buy independence through separate development [Avizienis, 1985]. Our case differs in one structural respect that makes it worse. In multiversion programming the shared artifact is a specification, and versions can at least diverge in how they read it. Here the shared artifact is an executable set of constants imported by both paths, so there is no input on which they could disagree about a status code. The correlation is not statistical but total.

5 What did catch it

The errors were found on a second pass over the same data, undertaken because the results felt thin rather than because anything had flagged them. Three checks did the work, and none compares two derivations of a quantity. Each goes back to the source’s own account of itself.

Read the field documentation, not the data. The vocabulary had been inferred from the values that appeared in the data, which is how docking and attachment came to be read as statements about where an object is. GCAT publishes a page defining every status value. It says that a record represents a phase in an object’s flight history and that the status names the event that ends the phase. Reading that separated “the object is somewhere” from “this record stops here.”

Decompose every unexplained residue. The first write-up reported 769 of its disagreements as unexplained. Stating a residue feels like rigour, which is why errors hide there comfortably. Asking

what the 769 consisted of ended most of the finding: 640 paired a CelesTrak impact or landing with a GCAT transition code, 500 of them docking or attachment. Gemini 8 is the clearest case. GCAT ends the record at docking with the Agena target vehicle and CelesTrak records the landing, and both are right about different events in one mission.

Ask whether the source already models what you accuse it of omitting. The tracking-loss finding rested on an implicit claim that CelesTrak has no way to record lost tracking. Checking that claim directly, rather than the count derived from it, found the field in use.

Redundancy asks whether two routes from A to B agree. These three ask whether A is what we think it is.

6 The correction was wrong too

The first correction moved C (collided) into the destroyed set, yet GCAT’s definitions allow an object to survive a collision. Checking that one entry against the source found more than the entry itself.

The source’s history of an object is in a file we never read. GCAT’s definitions say a collision destroys the object “in which case there is no subsequent phase”, while a survivor starts a new phase; and that after E (exploded) the “next phase of this object is a debris fragment.” So E marks a boundary, not a fate, and C is a fate only when nothing follows it. Whether anything follows cannot be seen in the files the pipeline read, because each holds exactly one record per object. GCAT’s own index says where the rest is: an event catalogue described as holding “later phases in history for objects in primary catalogs”, listed on the same page as the four files we fetched, under a different heading.

Checking the correction against those histories. We appended each object’s event-catalogue phases to its record, took the latest phase as GCAT’s current account of the object, resolved docked and attached objects through the object they joined, and recomputed the comparison (Appendix B). Of the 23 objects coded C that both catalogues carry, 21 have a later phase; of the 78 coded E , 66 do. The count falls from 261 to 220. It loses 42 objects, all 30 explosions and 12 of the 14 collisions in the 261. For 40 of them GCAT records a later phase that is still in orbit, which is where CelesTrak has them; among these are Cosmos 2251 and Iridium 33 after their 2009 collision and Fengyun-1C after the 2007 anti-satellite test. The other two are explosions with no later phase recorded, on which GCAT makes no present claim. One object enters, a payload whose later phase ends in a deorbit that CelesTrak does not record.

A second reading from the source agrees. GCAT publishes a derived catalogue of the most recent phase of every object. Read against it, the count is 219, and 208 objects are common to both readings. Every one of the 23 differences has a named cause: 12 are reentries recorded by one register and not yet by the other in the six weeks between retrievals, 3 are GCAT record revisions in that window, and 8 are deep-space objects the pipeline sets aside by design and the derived catalogue follows. Both readings drop all 42 objects. We treat 220 as the reference count, with the caveat that its event-catalogue input was retrieved on 29 September (GCAT’s data update of 28 September) rather than 18 August.

What this shows. The first correction was made by reading the documentation, as the paper recommends, and it still misread two codes and missed the file that would have settled them. It was a new reading of the vocabulary, and the pipeline’s gate and our three checks both inherited it. None of them flagged the 42 objects (Section 7); a later line-by-line reading of the source’s definitions did.

7 Three gates, implemented and measured

The checks in Section 5 were done by hand, and a discipline that depends on someone deciding to look again is not a discipline. We implemented all three and ran each on the defective code and on the first correction. Constraint-based data validation checks declared expectations about data at scale [Schelter et al., 2018]; these gates instead check the pipeline’s reading of a vocabulary, which

Table 2: Each gate on the defective code, and on the first correction, which still contained the second error of Section 6 (42 objects).

Gate	Defective code	First correction, against the second error
Total-function check on the status column	Fires: 9 unclassified codes over 1,099 rows (ATT, C, DK, E, EVA DP, GRP, N, REL, TFR)	Silent: 0 unclassified codes. E and C are classified, wrongly
Residue decomposition	Shows it: 20 groups, top 5 hold 729 of 932; 500 pair docking or attachment with impact or landing	Shows it, unread: 44 of 261 in groups pairing E or C with a CelesTrak object in orbit
Source inventory reconciliation	Fires: satcat100k listed, never fetched	Silent: scoped to the four main catalogues; the event catalogue is on the same index

is where the defect lived. The second error gives a test the first did not: the gates were built after the first error was known, but the second was unknown when they were built, so how they behave on it is closer to a prospective measurement. Table 2 gives the result.

The total-function check tests completeness, not meaning. Requiring every observed status to be assigned to a named category attacks the default that did the damage in Section 4, and on the defective code it lists every code implicated in the error. It cannot tell a right assignment from a wrong one. The first correction assigned E and C explicitly, to the wrong category, and the check passed. It also has a scoping cost: applied to every controlled column rather than the one the classifier consumes, it fires identically on both versions on 70,292 rows of CelesTrak’s orbit type, 70,130 of object type and 1,292 of data status, columns the pipeline treats as opaque strings. The scope need not be chosen by hand. The names both gate paths import from one module are exactly the definitions the gate cannot test (Section 4); at the published commit the vocabularies among them are all classifications of the GCAT status column, which is where the check separates cleanly.

Residue decomposition is an auditing aid, not a verdict. Grouping the unexplained disagreements by the two registers’ fields made the first error legible, since a status meaning “joined another object” should not systematically meet a landing in the other register. It made the second error legible too: on the corrected code, 30 explosions and 14 collisions sit against CelesTrak objects in orbit. We read that table after the correction and saw confirmation. The decomposition has no rejection criterion of its own, and the reader supplied the wrong one.

Inventory reconciliation is exactly as good as its list. Comparing the files the source lists against the files the harvest loaded found the missing fourth catalogue at once, with no false positive. On the second error it was silent, because we had scoped it to the section of the index headed as the main catalogues, while the index also lists the event catalogue, and eleven object catalogues in all. Widening the list would have caught it and would also flag files the pipeline has no reason to read. Running this gate also corrected our own build report, which said the missing file recovered 280 objects; it recovers 278, forced by the 900 and 622 that both reproduce exactly. A pass undertaken to correct three published numbers introduced a fourth error.

A check that looked independent and was not. The SHACL validation returned 3,564 results, exactly the sum of six of the seven counts. At the time this read as corroboration. It is an arithmetic identity: the shapes select the nodes the counts count, from the graph the counts were computed over, so when the counts were wrong the total was wrong by the same amount and still matched. A suite of such steps produces a great deal of green while measuring one thing repeatedly.

8 Asking a model for an independent check

The pipeline was written by people. We tested one part of the question for automated authors: asked to add an independent second path for the disposition count, does a code-generating model produce one whose meaning comes from somewhere other than the path it was shown? Reuse is the expected answer, since not duplicating a definition is correct practice by every other standard, and generated

Table 3: Controlled replication. Each cell is the number of generated paths that return the published, defective 932 on the frozen data, out of the trials run. “+ definitions” appends GCAT’s status definitions verbatim. All 72 paths that return 932 count the same status pairings as the path they were shown, on the probe or, for one, on the real files.

Request	opus-5-5	sonnet-5	haiku-4-5
Neutral	5/5	5/5	5/5
“Independent” path	5/5	5/5	5/5
Sharing forbidden	4/5	4/5	5/5
“Independent” + definitions	5/5	5/5	5/5
Sharing forbidden + definitions	4/5	5/5	5/5

test oracles are known to encode the behaviour of the code in front of them rather than the behaviour intended [Konstantinou et al., 2024].

A preliminary run, and what was wrong with it. A preliminary run of 30 trials on 19 August used a coding-agent CLI with the aliases `sonnet` and `haiku` and labelled outputs by pattern matching over the generated source. It had four problems. The aliases were not resolved to model identifiers in any record we kept. The CLI’s default tools were not disabled, so we cannot show that no trial read files. The prompt paraphrased the consumer function: it counted only explicit membership of the gone and in-orbit sets, where the real code treated every other status as in orbit. On the frozen data that paraphrase returns 208, not 932, so the trials never saw the defect they were scored against. And the loader bodies were elided, so no trial could see GCAT’s column names. What survives is narrower than we claimed. None of the 30 outputs introduced a status code absent from the prompt, and executed on the real files, 22 of them return exactly the 208 of the path they were shown, including all 19 that ran under the neutral and “independent” requests. When code sharing was forbidden, three of ten reproduced 208 and seven diverged, none by consulting the source: four guessed a column name and return 0, one parses differently and returns 189, one crashes, and one reads GCAT’s phase end date as a decay date and returns 795 (Appendix C).

A controlled replication. We reran the experiment with the pre-correction code verbatim, loaders included, so that the path shown returns the published 932. Each model was pinned by identifier and called through the same CLI with tools disabled, a one-line system prompt and an empty working directory: `claude-opus-5-5`, `claude-sonnet-5` and `claude-haiku-4-5-20251001`. Beside the three original requests, two new conditions append GCAT’s own status definitions verbatim, which is the remedy this paper recommends. Outcomes are behavioural. Each generated path is executed on a synthetic catalogue in which each pairing of a status code with a CelesTrak state holds a distinct power of two of objects, so any count decodes to exactly which pairings the path counted, and then on the frozen real files (Appendix C).

Table 3 gives the result. Of 75 trials, 72 return 932. Of the other three, one `claude-sonnet-5` path returns 0 because its hand-written file parser fails, and two `claude-opus-5-5` paths refuse to produce a count. Only one of the 75 would have exposed the defect: with code sharing forbidden, an Opus path wrote its own completeness check, and on the real files it stops and lists the unclassified codes, among them E, C, DK and ATT. That is the first gate of Section 7, produced unprompted, once.

The definitions did not change what was computed. Of the 30 paths written with GCAT’s definitions in the prompt, 29 return 932, and every one of the 30 counts E, C, DK, ATT, TFR and GRP exactly as the shown path does wherever the probe can see it. The paths read the documentation and kept the comparison. One Opus path annotated E as “exploded; the object continues as a debris fragment” and C as “collided; may or may not survive, so not treated as gone”, both correct, then placed both, with the docking codes, in a class of objects “still in space or transition” and compared that class with CelesTrak’s decay flag as the defective path does. On the real files its own completeness check stops it on the undocumented suffixes L? and R?; on the probe, which has none, it counts exactly what the defective path counts. A `claude-sonnet-5` path’s docstring says it “re-derives the GCAT status taxonomy straight from the phase documentation”; it returns 932.

What this does and does not establish. On this task, asking for an independent path, forbidding shared code, and supplying the source’s definitions each produced paths that, where they worked,

computed what the defective path computes, while looking to a code reviewer like separate implementations with new names and comments citing the documentation. It does not establish a rate for models in general (Section 9). If a check is meant to be independent, independence has to be specified as a property of its inputs and verified by what the check computes, not requested; models also struggle to correct their own reasoning without external feedback [Huang et al., 2024].

9 Limitations

This is one incident in one pipeline, with one family of defects. We can show that a redundancy gate failed this way here and name the structure that made it fail, which is two paths importing one set of definitions, but we make no claim about how often that structure occurs or how often it fails elsewhere. A rate would need a corpus of pipelines with known ground truth, which is the scarce thing. The same bound applies to the three gates. They are measured on one codebase against two errors, and the second error, found after the gates existed and not by them, is a single prospective test, not an evaluation.

The reference count is not an oracle. Neither register is treated as ground truth, and 220 is the number of objects on which the two make incompatible claims when GCAT is read through its own phase histories. Those histories come from an event catalogue retrieved six weeks after the other files, filtered to phases that began before the freeze; corrections GCAT made to older phases in that window cannot be separated out. A second reading from GCAT’s derived catalogue gives 219 and agrees on 208 objects, with every difference assigned a cause, and 8 of the differences are deep-space objects the comparison excludes by design. Two explosions with no recorded later phase are set aside; counting them as destroyed, as the first correction did, gives 222. The first correction was presented as final, and a further pass could revise 220 as well.

The generation experiment is small: one task, three models from one family, five trials per cell. A second family was planned and dropped, so nothing here speaks to models outside it. No trial could retrieve anything: the definitions were supplied in the prompt, and the event catalogue that settles E and C was supplied to none, so the experiment does not test an agent that can fetch the source for itself. We add one further limit: we are the party that benefits from the results being true, which is a reason to weight the released prompts, raw outputs and behavioural probe over our description of them.

10 Conclusion

A verification gate built on redundant computation reported agreement on seven published counts, three of which were wrong, one overstated more than fourfold. The gate was not defective. Its two paths differed in language, data structure and representation, and they agreed because both imported one set of constants encoding one misreading of a source catalogue’s status vocabulary. Diversity below a shared definition cannot detect an error in that definition, and the more elaborate the redundancy, the more convincing the resulting green.

The case supports a design rule rather than a general law. The definitions two gate paths share can be listed mechanically, and each of them needs a check that goes back to the source rather than to another derivation. That rule is necessary but not sufficient, as our own correction showed: it was made by reading the documentation and still misread two codes, and the checks we built to prevent a recurrence were silent on it. A further, line-by-line reading of the source’s definitions found it. Here redundant computation stood in only for a check of implementation, and the errors that reached publication, twice, were errors of meaning.

Acknowledgments and Disclosure of Funding

We thank the anonymous reviewers and the area chair for their comments, and Jonathan McDowell for maintaining GCAT and documenting it as thoroughly as he does.

References

- Algirdas Avizienis. The N-version approach to fault-tolerant software. *IEEE Transactions on Software Engineering*, SE-11(12):1491–1501, 1985. doi: 10.1109/TSE.1985.231893.
- CelesTrak. SATCAT: Satellite catalog, 2026. Retrieved 18 August 2026. <https://celestrak.org/pub/satcat.csv>.
- Bei Chen, Fengji Zhang, Anh Nguyen, Daoguang Zan, Zeqi Lin, Jian-Guang Lou, and Weizhu Chen. CodeT: Code generation with generated tests. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2207.10397.
- Dave E. Eckhardt and Larry D. Lee. A theoretical basis for the analysis of multiversion software subject to coincident errors. *IEEE Transactions on Software Engineering*, SE-11(12):1511–1517, 1985. doi: 10.1109/TSE.1985.231895.
- Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. Large language models cannot self-correct reasoning yet. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.01798.
- John C. Knight and Nancy G. Leveson. An experimental evaluation of the assumption of independence in multiversion programming. *IEEE Transactions on Software Engineering*, SE-12(1): 96–109, 1986. doi: 10.1109/TSE.1986.6312924.
- Holger Knublauch and Dimitris Kontokostas. Shapes constraint language (SHACL). W3C Recommendation, July 2017. Published 20 July 2017. <https://www.w3.org/TR/shacl/>.
- Michael Konstantinou, Renzo Degiovanni, and Mike Papadakis. Do LLMs generate test oracles that capture the actual or the expected program behaviour?, 2024. arXiv:2410.21136.
- Bev Littlewood and Douglas R. Miller. Conceptual modeling of coincident failures in multiversion software. *IEEE Transactions on Software Engineering*, 15(12):1596–1614, 1989. doi: 10.1109/32.58771.
- Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI scientist: Towards fully automated open-ended scientific discovery, 2024. arXiv:2408.06292.
- Jonathan C. McDowell. General catalog of artificial space objects, 2020. Retrieved 18 August 2026. <https://planet4589.org/space/gcat>. CC BY 4.0.
- Sebastian Schelter, Dustin Lange, Philipp Schmidt, Meltem Celikel, Felix Biessmann, and Andreas Grafberger. Automating large-scale data quality verification. *Proceedings of the VLDB Endowment*, 11(12):1781–1794, 2018. doi: 10.14778/3229863.3229867.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2203.11171.

A Object-level ledger of the disposition count

The unit throughout is an object present in both CelesTrak’s catalogue and GCAT’s main catalogue, keyed by catalogue number (69,390 objects). Two distinctions of unit matter when reading the figures below: 998 counted objects removed from the comparison, not disagreements, and 500 is the docking and attachment subset of the 769, not all of it. Table 4 tracks every object through the three versions; `count_ledger.py` produces each figure and asserts that the arithmetic closes.

Table 4: The disposition count, object by object. Status groups refer to the object’s GCAT status code.

Published count	932
GCAT gone, CelesTrak gives no decay date	163
CelesTrak decayed, GCAT not gone	769
docking or attachment (DK, ATT)	500
other transition codes (TFR, GRP, N, ...)	140
explosion or collision (E, C)	57
in orbit but tracking lost (OX)	31
error entries (ERR)	21
in orbit (O, AO)	14
gone, but status carries a trailing ?	6
Left the comparison at the first correction	−724
transition codes now excluded (of 998 objects excluded in all)	−640
explosions and collisions reclassified as destroyed, CelesTrak agrees	−57
error entries now excluded	−21
trailing ? now stripped	−6
Entered at the first correction	+53
explosions and collisions reclassified as destroyed, CelesTrak silent	+44
trailing ? now stripped	+9
First correction	261
Left at the full-history check: later phase in orbit or unrecorded	−42
explosions	−30
collisions	−12
Entered: attached object whose later phase ends in a deorbit	+1
Full-history count	220

B The full-history check

`history_check.py` appends to each object’s record the phases GCAT’s event catalogue lists for the same identifier, drops the 19 phases that began after 18 August 2026, and takes the status of the latest phase. Following GCAT’s definitions, C counts as destruction only as a last phase; E, N, LEASE and REFLT as a last phase leave the present state unrecorded (28 objects, set aside); docking, grappling, attachment and transfer codes resolve through the object named as destination. The three catalogue numbers GCAT assigns to two objects each are resolved on both, and count only if both agree. The 14 objects whose last phase is a lease are all in orbit in CelesTrak, so treating a lease as a free-flying state instead leaves the count unchanged.

`currentcat_crosscheck.py` repeats the comparison from GCAT’s derived catalogue of the most recent phase of each object, mapping every status text that occurs to a disposition and failing on any it cannot map. Over 69,249 comparable objects it gives 219, with 208 in common with the 220. The 12 objects only in our reading are 9 reentries CelesTrak recorded between 13 and 18 August that the frozen GCAT file did not yet carry, 2 collisions with no later phase that the derived catalogue now labels as explosions, and 1 record GCAT revised after the freeze. The 11 only in the derived reading are 8 deep-space objects, which the pipeline sets aside, and 3 reentries GCAT recorded after the freeze. Of the 42 objects removed from the first correction, the derived catalogue lists 40 as in Earth orbit and 2 as exploded; CelesTrak gives none of the 42 a decay date.

C Generation experiment: protocol and per-trial results

Behavioural probe. `behaviour_probe.py` writes a synthetic CelesTrak file and GCAT file under the paths the prompt names. There are 14 cells: each of E, C, DK, ATT, TFR and GRP, paired with a CelesTrak decay date and without one, plus two controls every reading counts (R with no decay date, 0 with one). Cell i holds 2^i objects, all with five-digit catalogue numbers so zero-padded and bare forms coincide. The generated code runs in a separate process with the prompt’s module importable and its names preloaded, since many paths use them without importing them; the original path is stubbed to return a sentinel that decodes to nothing. Every zero-argument function the code defines, other than loaders and helpers, is called, and any returned or printed count that decodes to

a set containing both controls is taken as the path’s signature. The same code is then run on the frozen real files and the number each chosen function returns is recorded. Reference signatures are obtained by running the reference code through the same harness, not written down; the verbatim pre-correction path returns 932 on the real files, which checks the harness end to end.

Preliminary run (19 August). Prompt in `reuse_experiment.py`; raw outputs in `reuse_raw/`; labels in `behaviour_labels.json`. The path shown counts only the two controls on the probe and returns 208 on the real files. Neutral request: 9 of 10 reproduce 208; 1 does not run (it imports a name the module does not define). “Independent”: 10 of 10 reproduce 208. Sharing forbidden: 3 reproduce 208; 4 read GCAT’s catalogue-number column under a guessed name and return 0; 1 parses the files differently and returns 189; 1 crashes on real rows; 1 ignores the status codes, treats any GCAT phase end date as a decay, and returns 795. The earlier regex labels called all 30 “inherits the defect”; that label described the vocabulary each path contains, not what it computes.

Replication (29 September). Runner `reuse_replication.py`; prompts, raw outputs and per-trial metadata in `replication_raw/`. Every model was called through the Claude Code CLI with `-tools ""`, `-system-prompt "You are a helpful assistant."`, `-setting-sources ""`, `-strict-mcp-config` and no session persistence, from an empty directory. Sonnet and Haiku ran on CLI version 2.1.277; `claude-opus-5-5` needs 2.1.280 or later and ran on 2.1.284. The CLI does not expose sampling parameters, and its defaults include extended thinking, which every trial used. It also makes a small fixed auxiliary call to a Haiku model; the recorded per-model token usage shows that in every trial the answer’s output tokens came from the requested model, in a single turn with no tool use. A local model from a second family was started and abandoned because it decoded too slowly to complete the design; no partial result is reported.

Per-trial outcomes. On the probe, 72 of the 75 paths have exactly the signature of the path shown, which counts all six codes as in orbit whenever CelesTrak records a decay. The three exceptions: one `claude-sonnet-5` path (sharing forbidden) returns 0 on probe and real files because its hand-written CSV parser fails; one `claude-sonnet-5` path (sharing forbidden, with definitions) takes its decay signal from CelesTrak’s operational status field, which the probe leaves blank, and returns 932 on the real files; one `claude-opus-5-5` path (sharing forbidden) raises on every status it has not classified, which on the real files includes ATT, C, DK, E, EVA DP and GRP. A second Opus path (sharing forbidden, with definitions) matches the shown path on the probe and on the real files stops on the undocumented L? and R?. Three `claude-haiku-4-5` paths and several Opus paths reference a constant of the prompt’s module without importing it or run as command-line programs; the harness supplies the constant and a data-directory argument, as the pipeline would, and reads counts that are printed rather than returned.

D Artifacts

Code, gate implementations, experiment prompts, raw model outputs, per-trial metadata and the behavioural probe are in the project repository at <https://github.com/fabio-rovai/space-object-register-ontology> under `paper/gates/`. The source files are public and were retrieved on the dates below; the SHA-256 prefixes identify the exact copies used.

File	Retrieved	SHA-256 (prefix)
CelesTrak <code>satcat.csv</code>	18 Aug 2026	888f67c53ee3d006
GCAT <code>satcat.tsv</code>	18 Aug 2026	33751fb8ac5b6fb1
GCAT <code>auxcat.tsv</code>	18 Aug 2026	82363eb712d8a392
GCAT <code>ftocat.tsv</code>	18 Aug 2026	d6b6351ceb0a5a98
GCAT <code>satcat100k.tsv</code>	18 Aug 2026	6ec65d36db9f6263
GCAT <code>ecat.tsv</code> (event catalogue)	29 Sep 2026	39023a287c27326e
GCAT <code>currentcat.tsv</code> (derived)	29 Sep 2026	8a0d9285b6fce7b1